

Reading Reference: Textbook 1 Chapter 8

Text 2: [Direct Execution](#)

EXCEPTIONS & HARDWARE SUPPORT

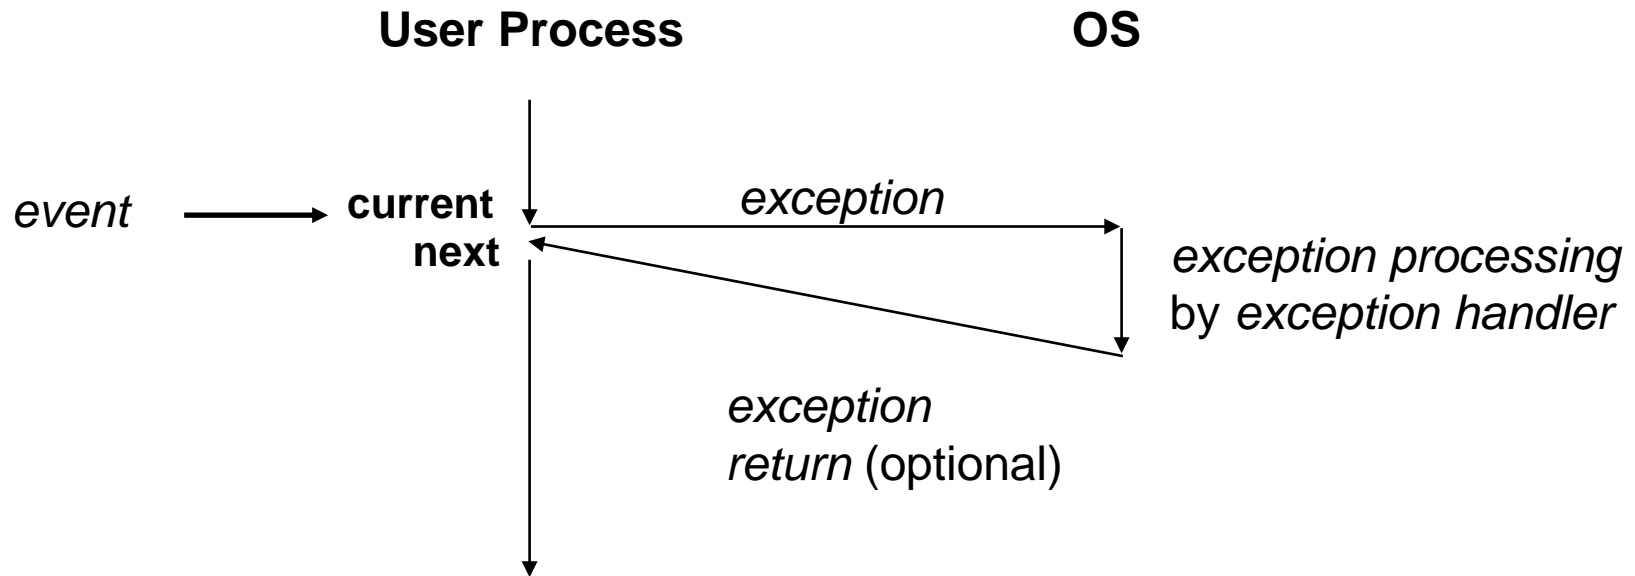
TANZIR AHMED
CSCE 313 Spring 2021

Altering the Control Flow

- CPU executes instructions 1-at-a-time mostly sequentially
- Program-assisted mechanisms for changing control flow:
 - *Jumps and branches*—react to changes in program state
 - *Function call and return using stack discipline*—react to program state
- Insufficient for a useful system
 - *The user application is the central thing – how to let OS into the CPU unless the app gives up control?*
 - *Thus, difficult for the CPU to react to other changes in system state*
 - Data arrives from a network adapter
 - Instruction divides by zero
 - User hits control-C at the keyboard
- System needs mechanisms for “exception control flow” ₂

Exception Control Flow

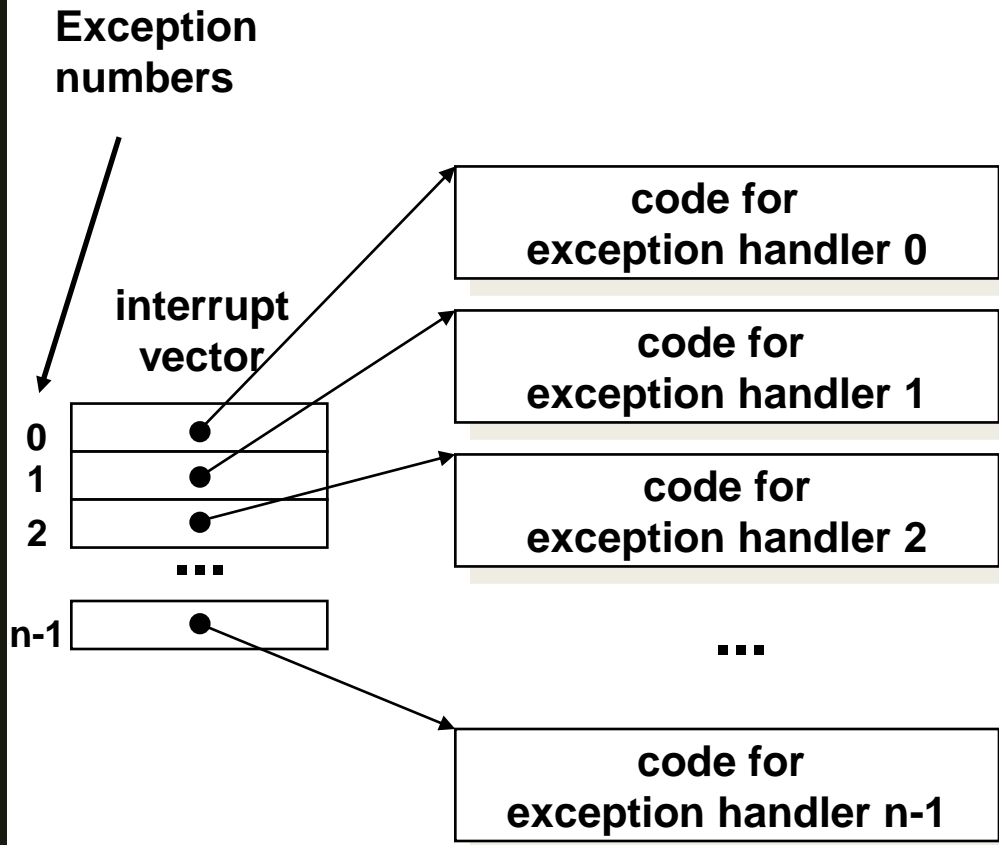
- An *exception* is a transfer of control to the OS in response to some event (i.e., change in processor state)
 - It also gives “asynchronous” behavior



Asynchronous Exceptions (Interrupts)

- Caused by events **external** to processor (i.e., outside the current program)
 - *Indicated by setting the processor's interrupt pin(s)*
 - *Handler returns to “next” instruction after servicing*
- **Examples:**
 - *I/O interrupts*
 - Key pressed on the keyboard
 - Arrival of packet from network, or disk
 - *Hard-reset interrupt*
 - Hitting reset button
 - *Soft-reset interrupt*
 - Hitting control-alt-delete to initiate restart on a PC

Interrupt Vectors




- Each type of event has a unique exception number k
- Index into jump table (a.k.a., interrupt vector)
- Jump table entry k points to a function (exception handler).
- Handler k is called each time exception k occurs.

Synchronous Exceptions: Traps, Faults, Aborts

- Caused as result of executing an instruction
 - *From within the CPU*
- 3 types:
 - *System Calls*
 - *Faults*
 - *Aborts*

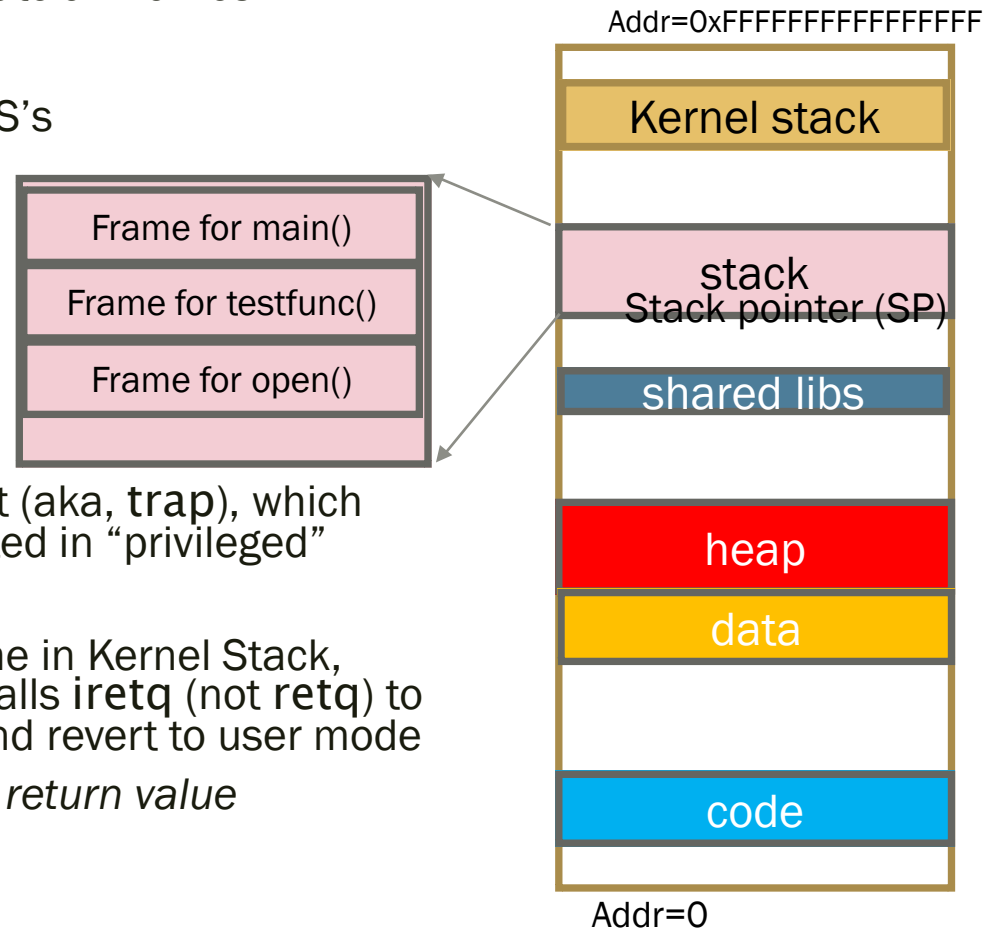
System Calls

- Functions called by user programs to interact with the OS or request core OS services
- However, unlike traditional functions, system calls must execute in some “privileged” mode so that has access to OS’s internal information and data structures etc.
 - *Hence, these are very different from regular function calls*
 - *Example: Opening a File*
 - *User calls `open(filename, options)`*
 - *OS must find (if reading) or create (when writing) the file, consult permissions and get it ready*
- However, when everything checks out, it must also act like a regular function that returns a file descriptor
- Let us see how a system call is implemented then

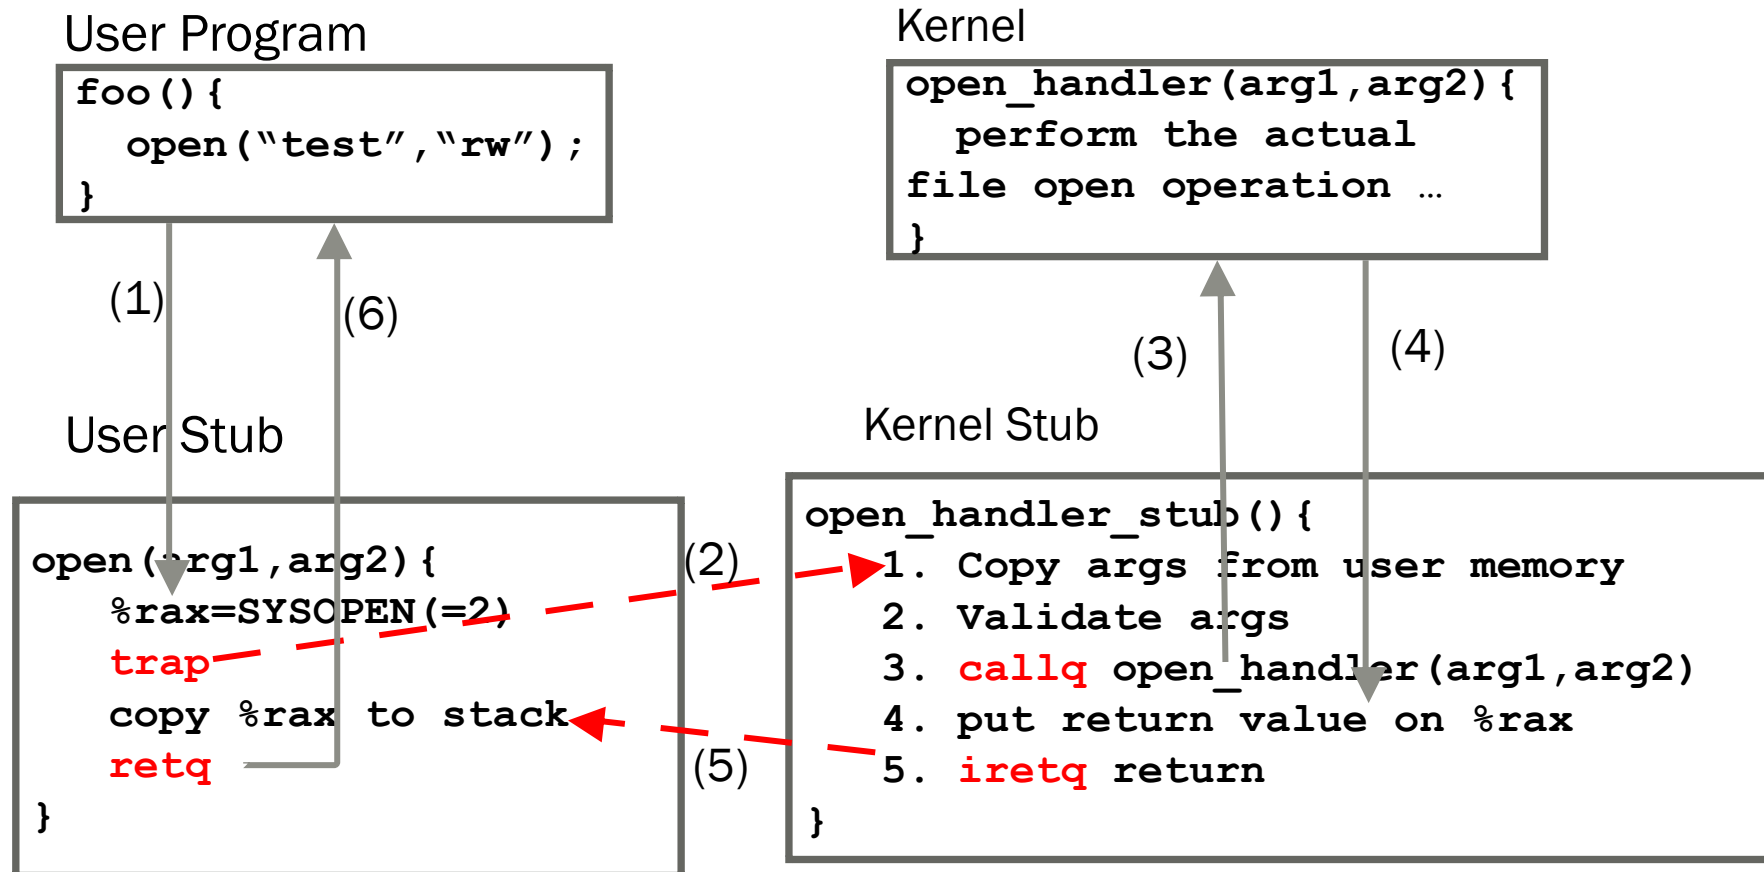


Mechanism of System Calls

- We need to take a closer look at the address space layout of program
- The user stack maintains a collection of stack frames, but this is insecure for syscalls because someone could manipulate the SP and navigate all frames
 - *E.g., gdb can do that*
- Therefore, the solution is to keep OS's internal data separately in a place called kernel stack
- The desired system call's opcode is placed in reg %rax (e.g., 0 for read, 1 for write, 2 for open, 3 for close)
- The CPU issues a software interrupt (aka, trap), which invokes an interrupt handler executed in "privileged" mode
- The interrupt handler makes a frame in Kernel Stack, places return value in %rax, then calls iretq (not retq) to return from the interrupt handler and revert to user mode
 - *%rax is then put in place of a return value*



Control Flow in System Calls

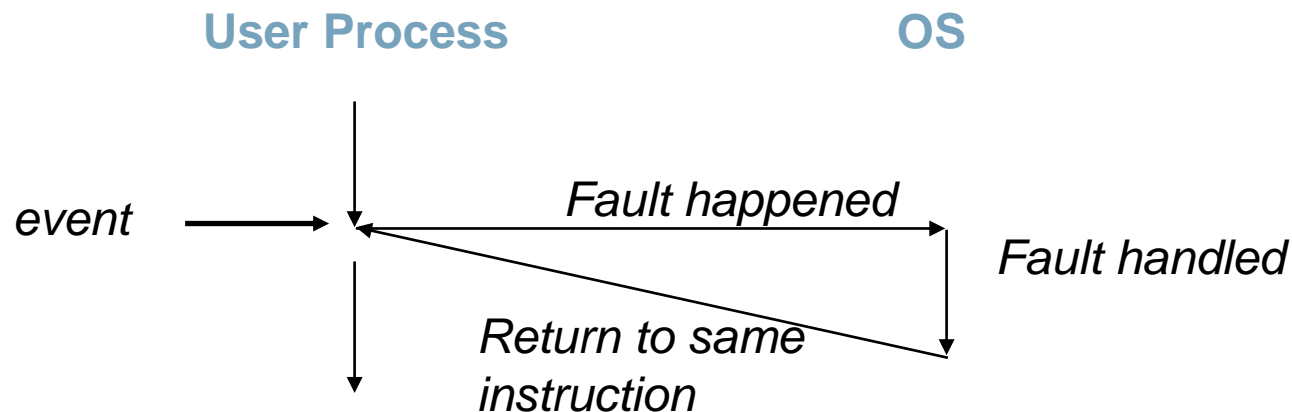


Dashed lines mean context switch due to software interrupt,
This is also where execution mode changes to/from "privileged"

Faults

■ Attributes

- *Unintentional but possibly recoverable*
- *Examples: Page Faults*
- *Either re-executes faulting (“current”) instruction or aborts*

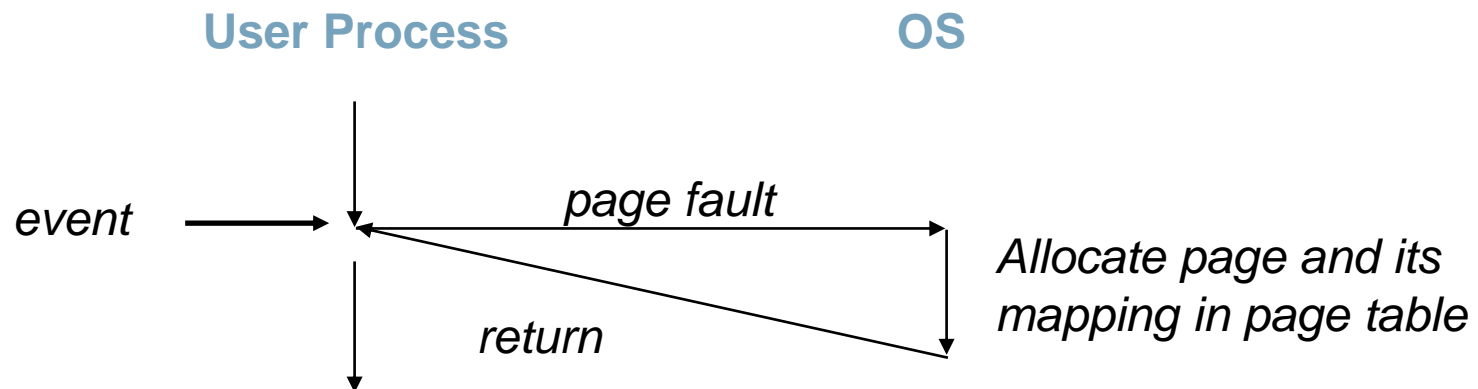


Fault Example #1

```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```

■ Memory Reference

- *User writes to memory location*
- *That portion (page) of user's memory is not mapped yet (because memory pages are mapped only when necessary)*
- *Page handler must load page into physical memory*
- *Returns to faulting instruction*
- *Successful on second try*




Fault Example #2

■ Illegal Memory Reference

- *User writes to memory location*
- *Address is not valid*
- *Page handler detects invalid address*
- *Sends `SIGSEGV` signal to user process*
- *User process exits with “segmentation fault”*

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```

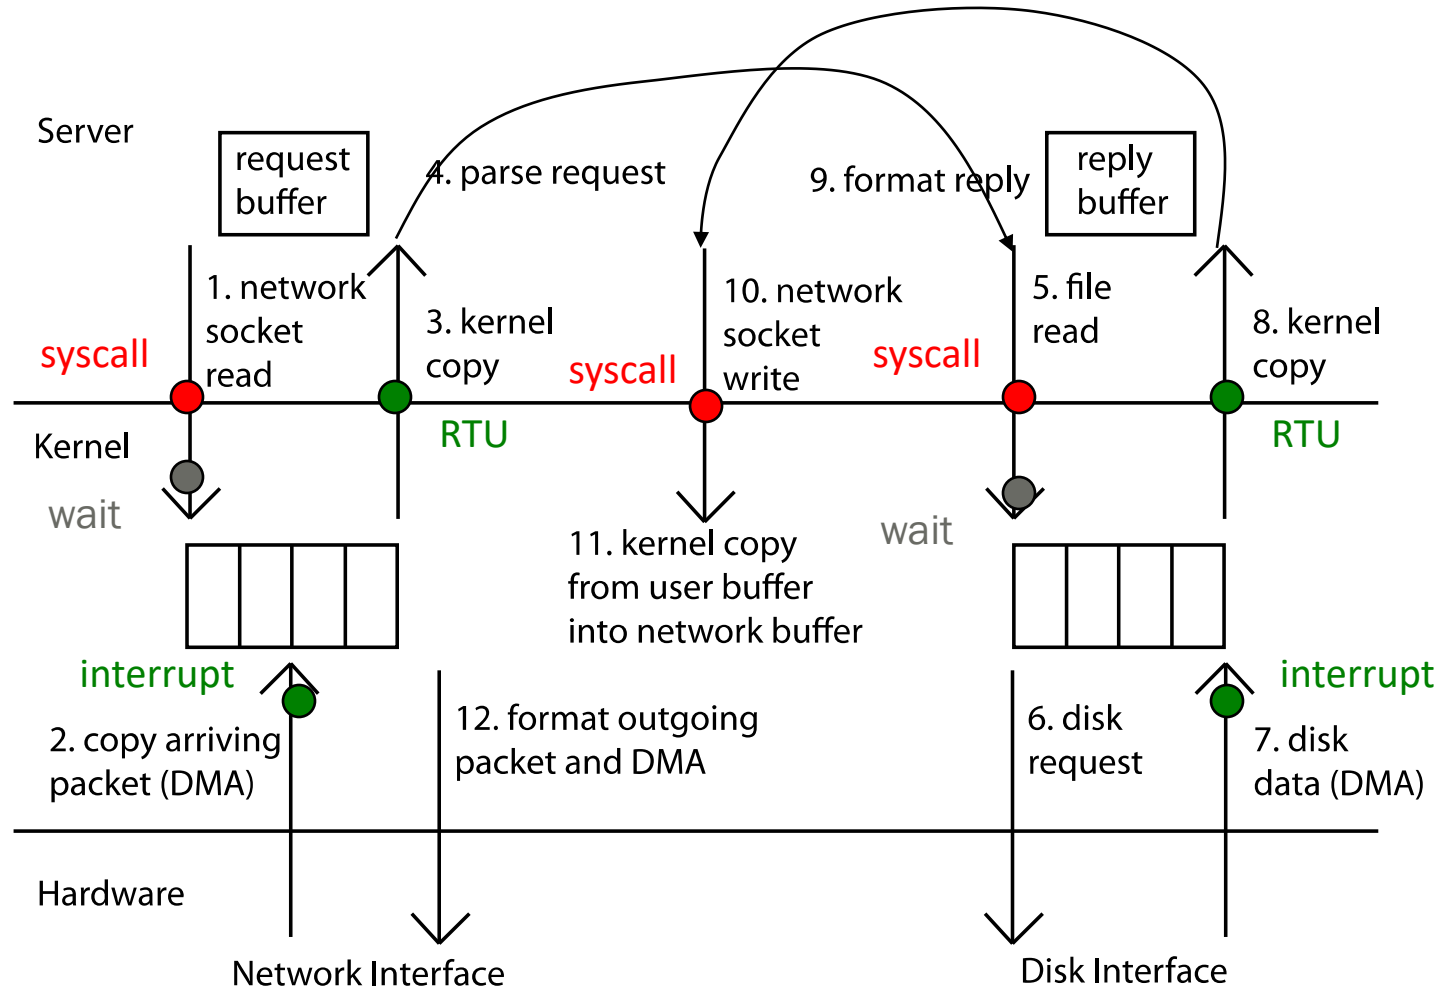


Aborts

■ Attributes

- *Unintentional and unrecoverable*
- *Examples: parity error, machine check, divide by zero*
- *Aborts current program and hands control over to the OS*
 - This is the way for the OS to put essential error checking
 - Also, an excellent way to make OS resilient when applications are failing or crashing

Example: Web Server




Summarizing Control Flow Exceptions

- User programs are not in charge of (and therefore not burdened with) handling every exception
 - *Not every program will check for divide-by-0 error*
- Mechanism is used by OS to do things beyond error handling
 - *E.g., page faults are used to enable “lazy” physical memory allocation*
- Are Synchronous/Internal (Traps, Faults, Aborts) OR Asynchronous/External (I/O Interrupts, Hard or Soft Reset etc.)

Architectural Support for OS

- A modern with all its features cannot run on a hardware that does not support it
 - *For example, the CPU must be Interrupt-ready if you want to run multiprogramming, time-sharing*

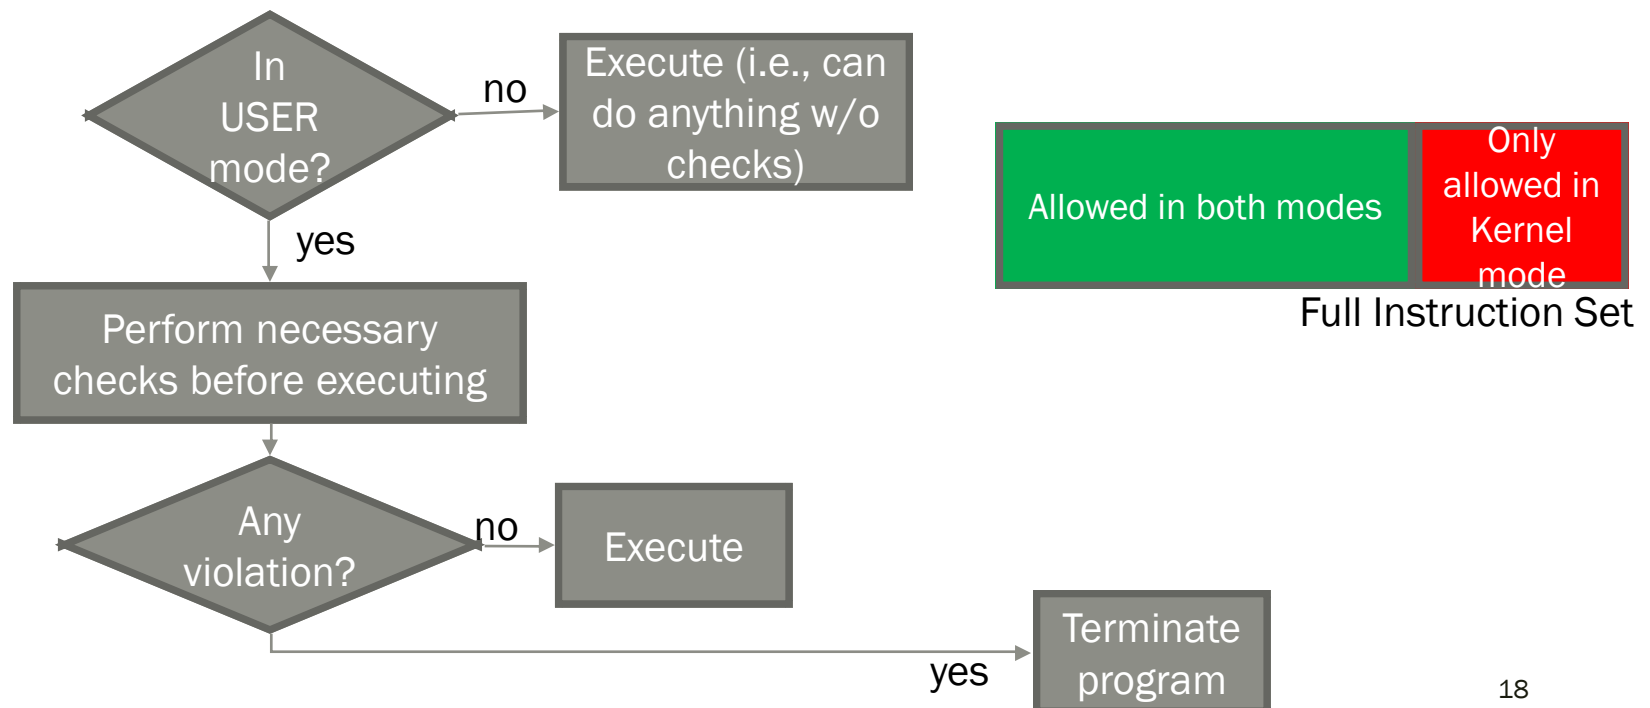


Architectural Features

- Protection Modes
 - *Protection Ring / 2 modes: User/kernel*
 - *Privileged Instructions*
- Interrupts and Exceptions
- System Calls
- Timers (clock)
- Memory Protection Mechanisms
- I/O Control and Operation
- Synchronization Primitives (e.g., atomic instructions)

Dual-mode Execution

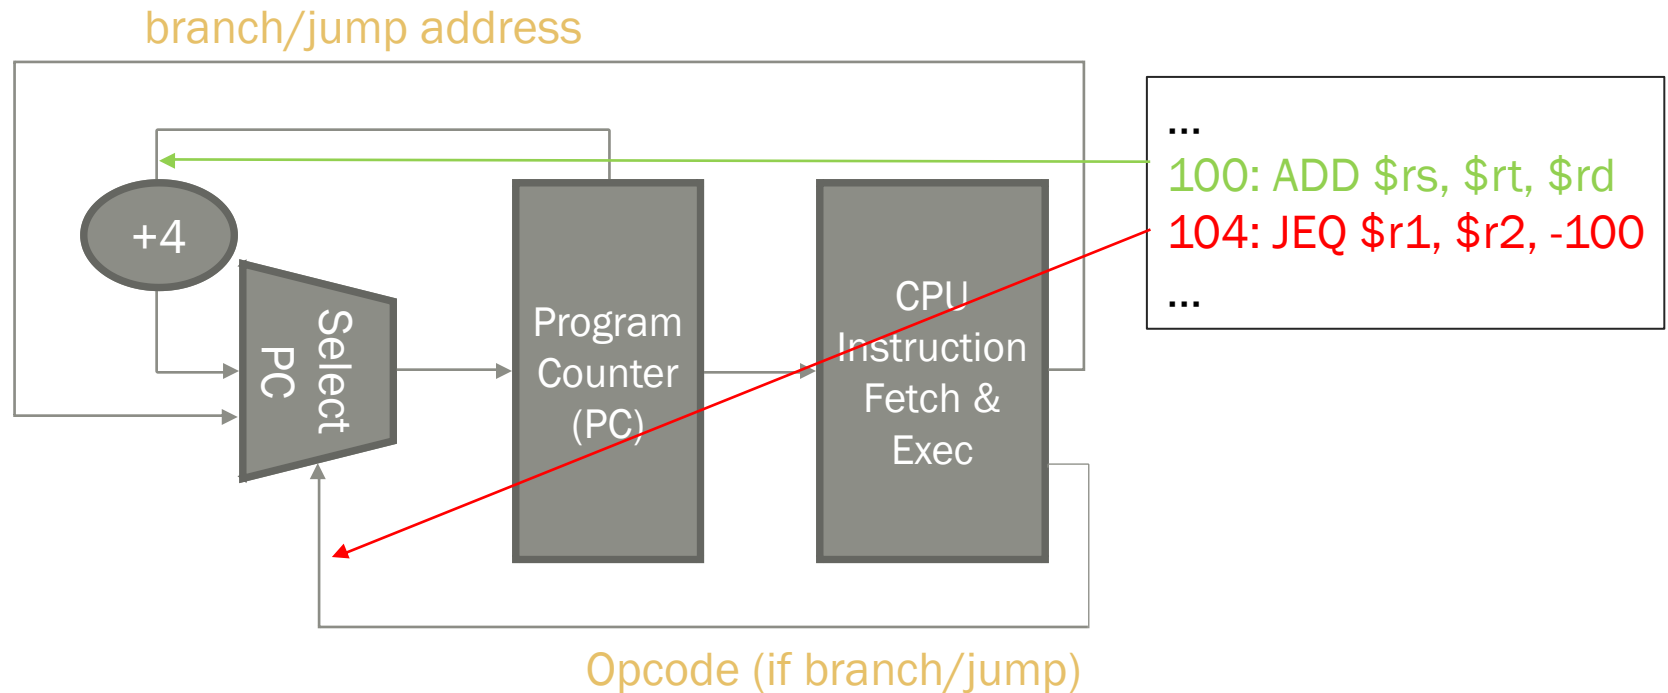
- Every CPU has at least 2 modes of execution (the CPU alternates between the modes)
 - **Kernel-mode:** Execution with the **full privileges** of the hardware
 - **User-mode:** Execution with **Limited privileges**
 - Only those granted by the operating system kernel



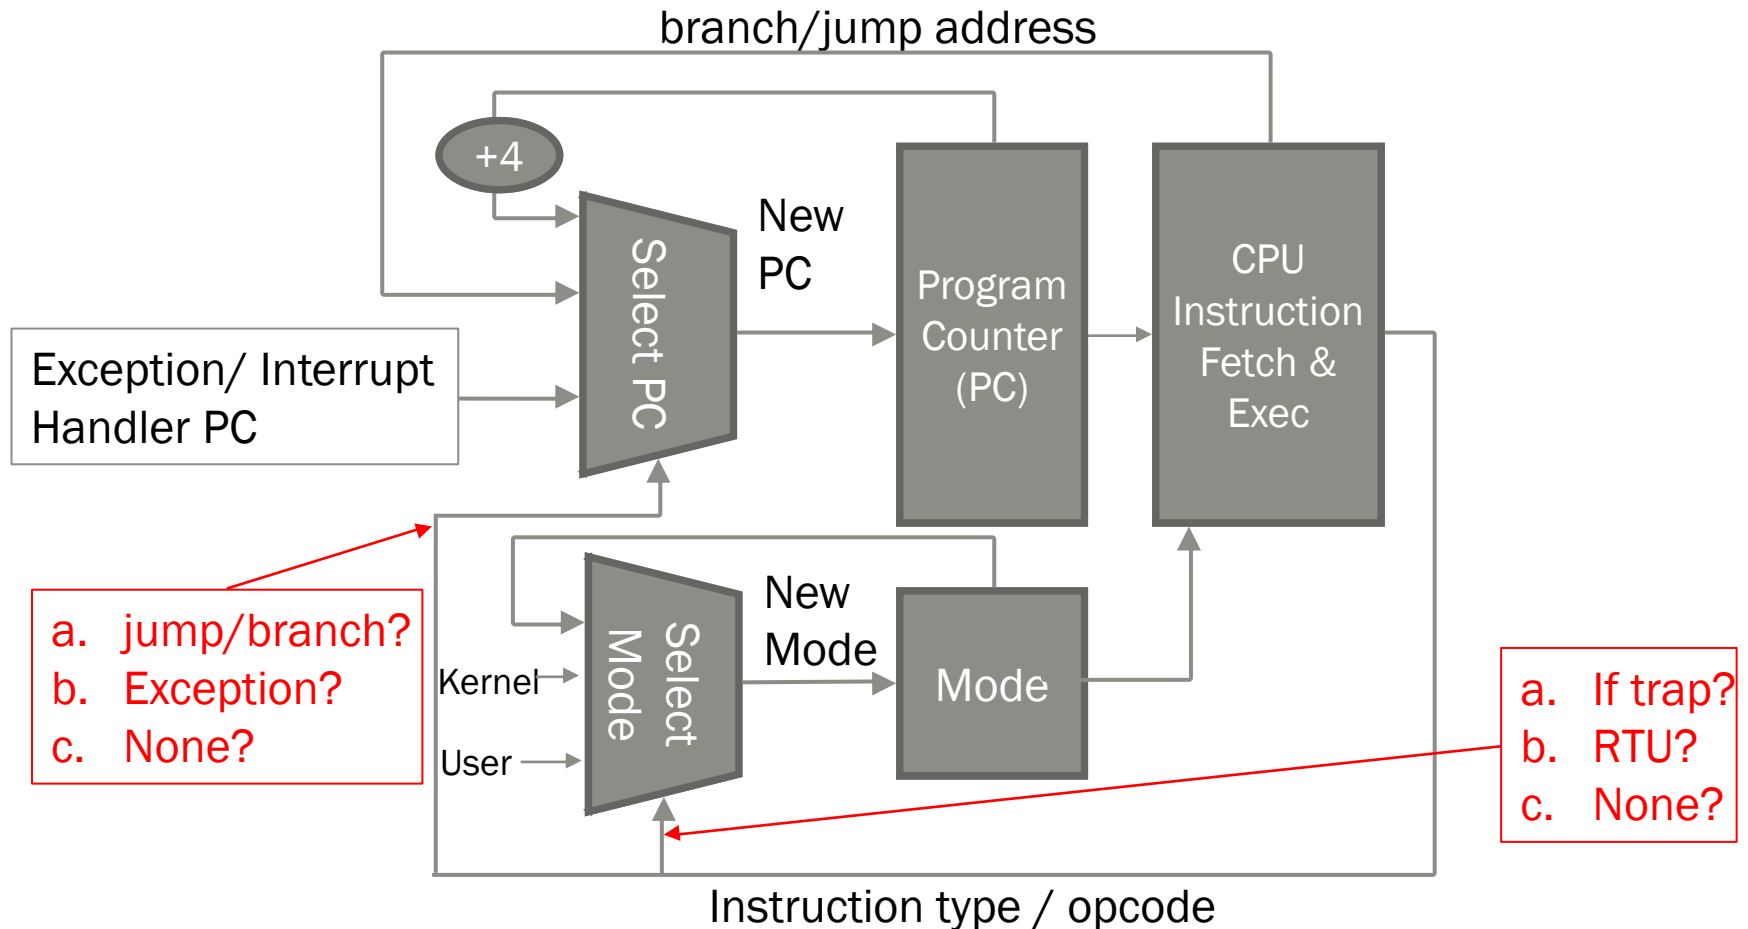
Privileged Instructions - Examples

- Only the OS should be able to
 - *Directly access I/O devices (disks, printers..)*
 - Allows OS to enforce security and fairness
 - User programs cannot possibly be fair to each other
 - *Manipulate memory management state*
 - E.g., page tables (Virtual-> Physical), protection bits, TLB entries, etc.
 - Processes use them, but cannot modify – that would defeat the protection
 - *Adjust protected control registers*
 - User $\leftarrow \rightarrow$ Kernel modes or Raise/Lower interrupt level
 - *Execute CLF instruction*

A Simple CPU Model – Needs Some Change



A CPU with Dual-Mode Operation



What are “Necessary” Checks?

- One example: Kernel Memory Protection

- Note that because of Virtual Memory, no process can harm another

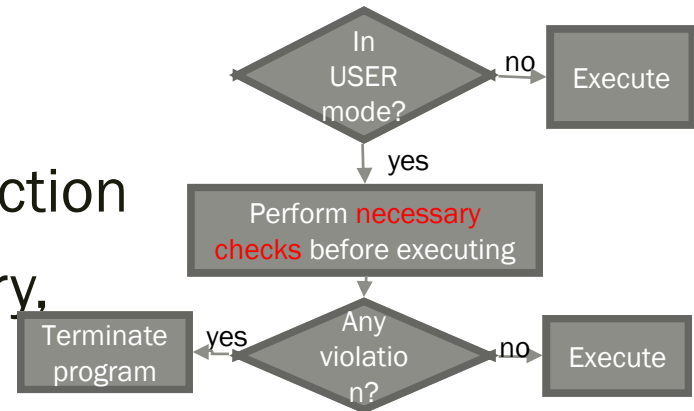
- *Then, why bother about memory protection?*

- However, the kernel is mapped into lower part of each process's address space like in the picture:


- Why is Kernel also included?

- *To load and run the process in the first place*
 - *Handle Interrupts, exceptions, and system calls efficiently w/o address space switch*

- Therefore, kernel must be protected from a faulty or malicious user program




Private address space



Memory Protection

- Could we just make Kernel memory read-only?
 - *No. Because kernel contain sensitive data that should not even be read*
 - *Example: page tables*
- A primitive but effective solution in hardware:
 - A **Base** and a **Bound** register for each process
 - Cannot access (read or write) anything below **base**
- This check is done only in User mode
 - *Kernel mode has unlimited access*



Hardware Timer

- Operating system timer is a critical building block
 - *Many resources are time-shared; e.g., CPU*
 - *Allows OS to prevent infinite loops*
- Fallback mechanism by which OS regains control
 - *When timer expires, generates an interrupt*
 - *Handled by kernel, which controls resumption context*
 - Basis for **OS scheduler**; more later...
 - *Setting (and clearing) a timer is a privileged instruction*

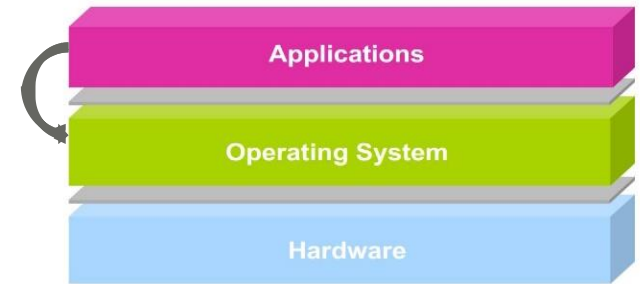
Mode Switch

- Mode separation does NOT mean that a user program cannot request a Kernel-mode operation
 - *User mode to kernel mode switch is very common*
 - *How do we switch from one mode to the other?*
 - Such that the protection is not compromised

User → Kernel Mode Switch


■ From user-mode to kernel-mode

- *Interrupts*
- *(Synchronous) Exceptions (Faults and Aborts)*
- *System calls (traps) (aka protected procedure call)*



Kernel → User Mode Switch

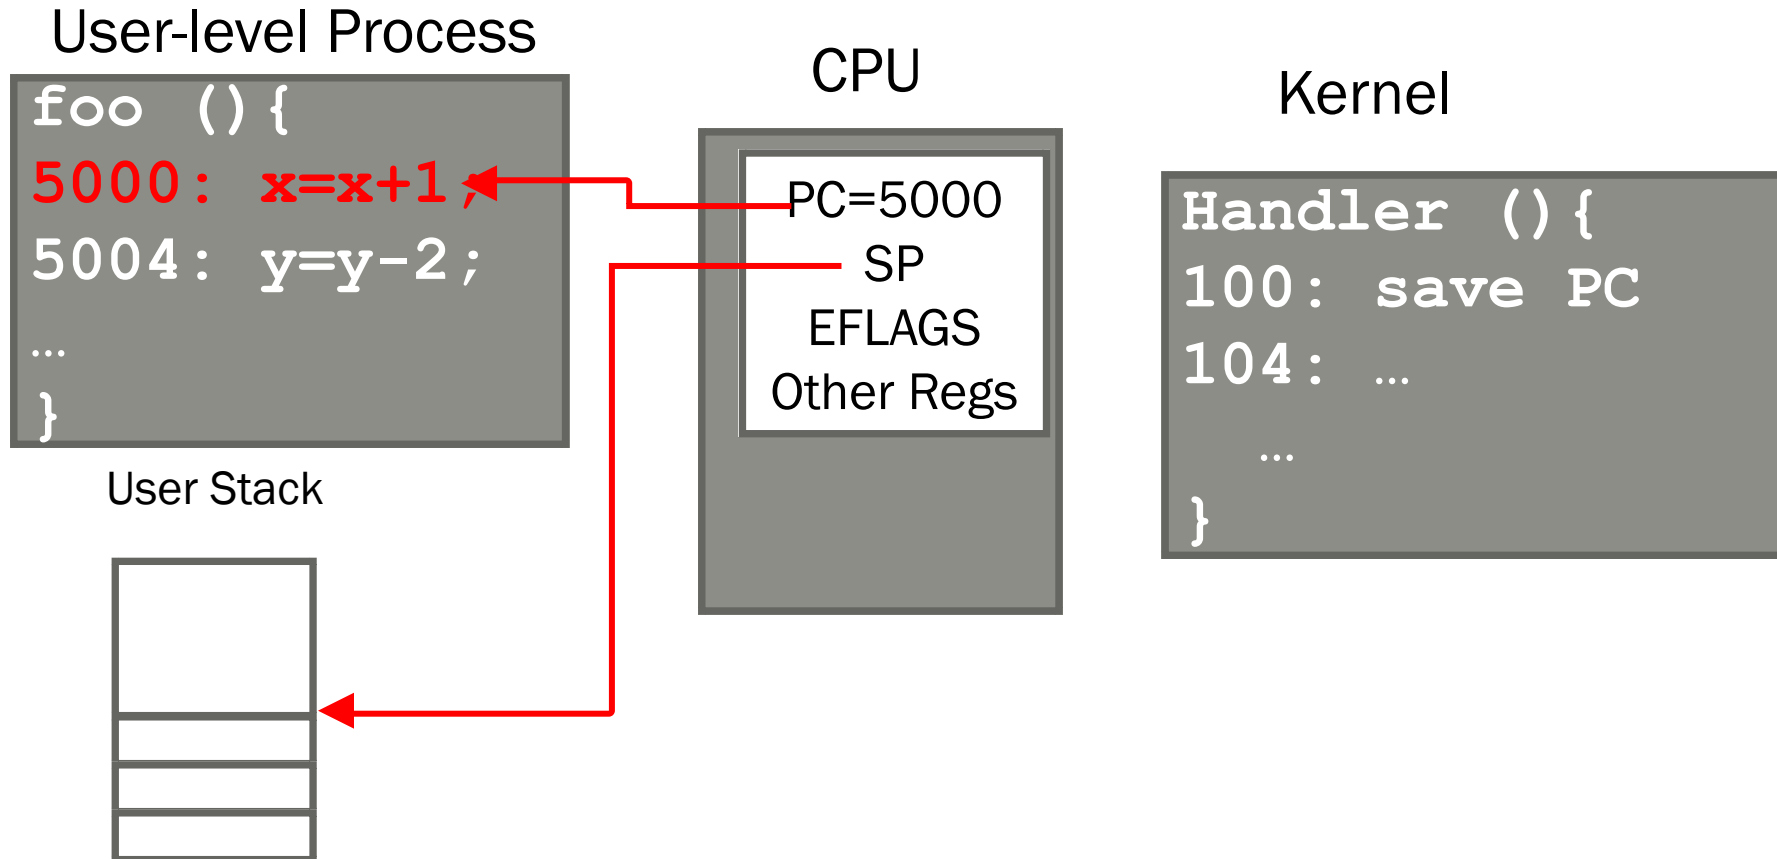
- From kernel-mode to user-mode
 - *New process/new thread start*
 - Jump to first instruction in program/th
 - *Return from interrupt, exception, system call*
 - Resume suspended execution
 - *Process/thread context switch*
 - Resume some other process



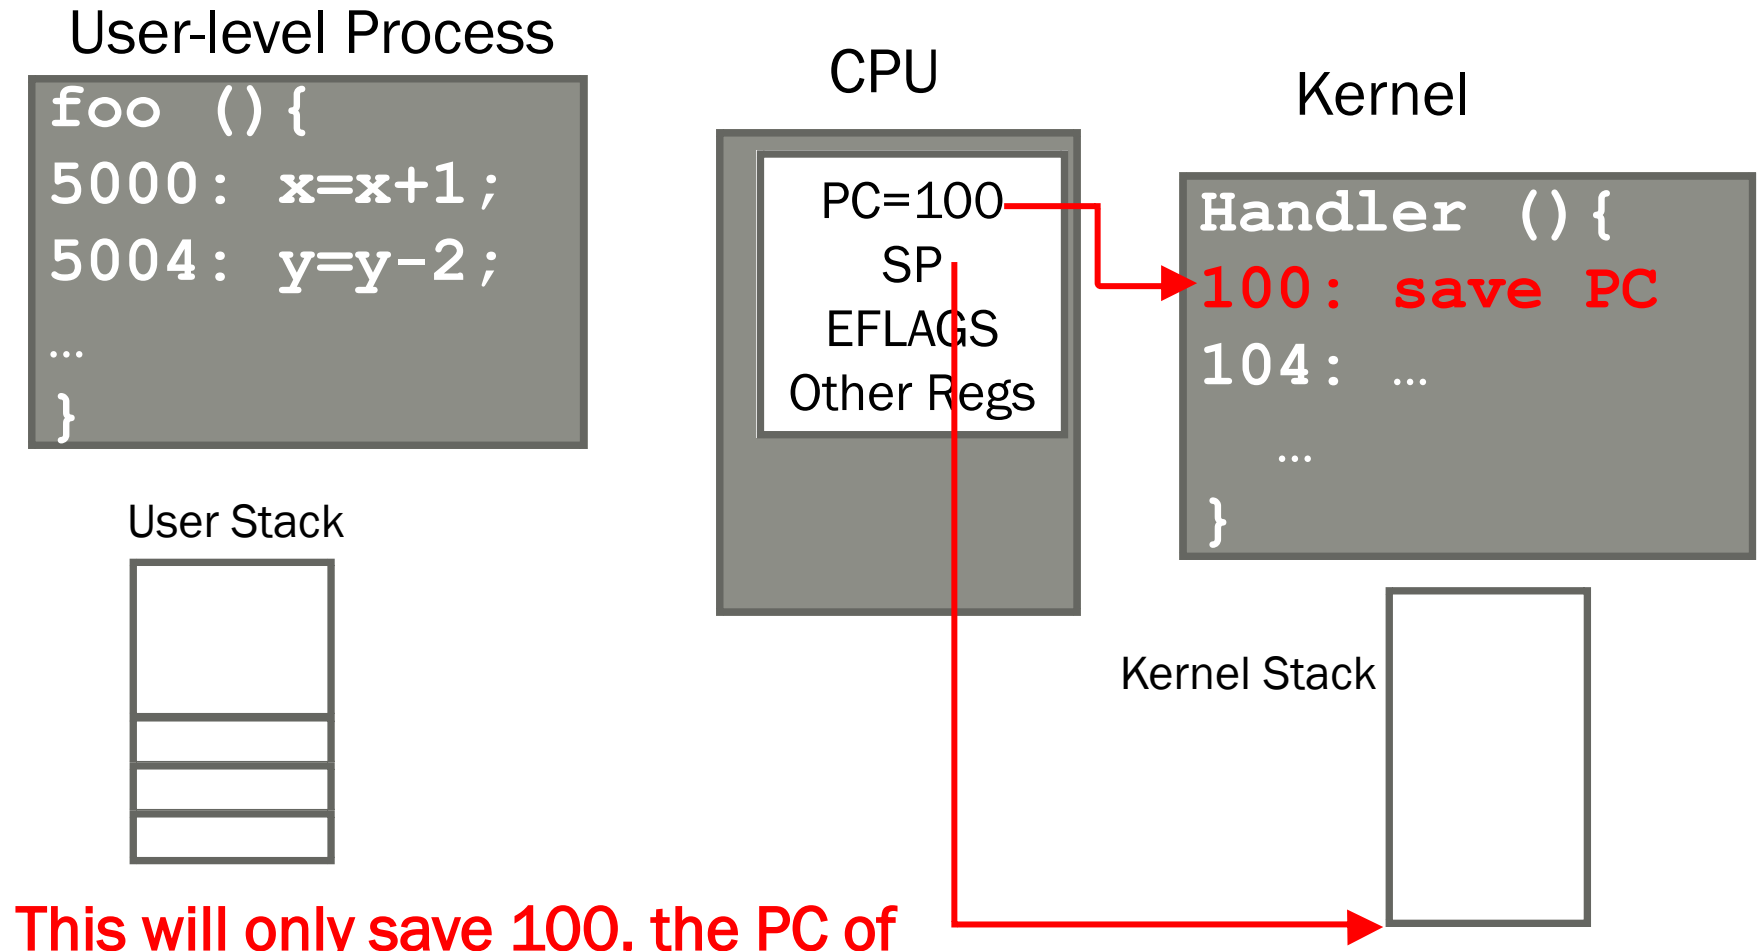
Safe Mode Transfer – Interrupt Handling

- First, User->Kernel. The main idea is rather simple
 - *Store the state of the running process (so that it can be resumed later)*
 - *Execute the handler*
 - *Return to the interrupted process by restoring the saved state*
- But the actual implementation is a bit more complicated
- We first need to know which process info must be stored
 - *Basically, the entire program/process state*

Saving Process State: Difficulty



Saving Process State: Difficulty



This will only save 100, the PC of the handler. The original PC=5000 of user app is lost forever!!

To Summarize


- The processor has **only 1 set of SP, PC, EFLAGS etc.**
- Any piece of code (e.g. handler code as well) will require its own PC (and also SP and others) first loaded into the CPU
- Switching from User code to Handler code means **overwriting** PC, SP etc. with the handler PC, SP etc.
 - *But ALAS!!! We just lost the PC, SP for the user code*
 - *How can we ever recover those??*
- Quoting the Anderson book: *“This is akin to rebuilding the car’s transmission while it barrels down the road 60mph”*
- Solution: Take hardware help
 - *Clearly, any other code will also need PC, SP,..*
 - *Hardware does not need to use SP, PC to implement a logic*

Plan of Action: Before Interrupt


User-level Process

```
foo () {  
  5000 : x = x + 1 ;  
  5004 : y = y - 2 ;  
}
```

User Stack




CPU



Kernel

```
handler () {  
  pushad  
  ...  
}
```

Kernel Interrupt Stack




Plan of Action: Hardware Action


User-level Process

```
foo () {  
  5000 : x=x+1 ;  
  5004 : y=y-2 ;  
}
```

User Stack



CPU



Kernel

```
handler () {  
  pushad  
  ...  
}
```




Kernel Interrupt Stack

Plan of Action: As Handler Starts


User-level Process

```
foo () {  
  5000 : x = x + 1 ;  
  5004 : y = y - 2 ;  
}
```

User Stack



CPU



Kernel

```
handler () {  
  pushad  
  ...  
}
```



Kernel Interrupt Stack

User to Interrupt Handler – Mechanism on x86

■ Hardware does the following:

1. *Mask further interrupts*
 - they are stored, not thrown away
 - If another interrupt arrives in steps 2-5 while the interrupted context is only partially saved, it is lost. Int. Disabled thus provides a “concurrency safety”
2. *Change mode to Kernel*
3. *Copy PC, SP, EFLAGS to the **Kernel Interrupt Stack (KIS)***
4. *Change SP: to the KIS (above the stored PC, SP, EFLAGS)*
5. *Change PC: Invoke the interrupt handler from the Interrupt Vector Table (i.e., overwrite PC with the handler PC)*

■ Software (i.e., the handler code) does the following:

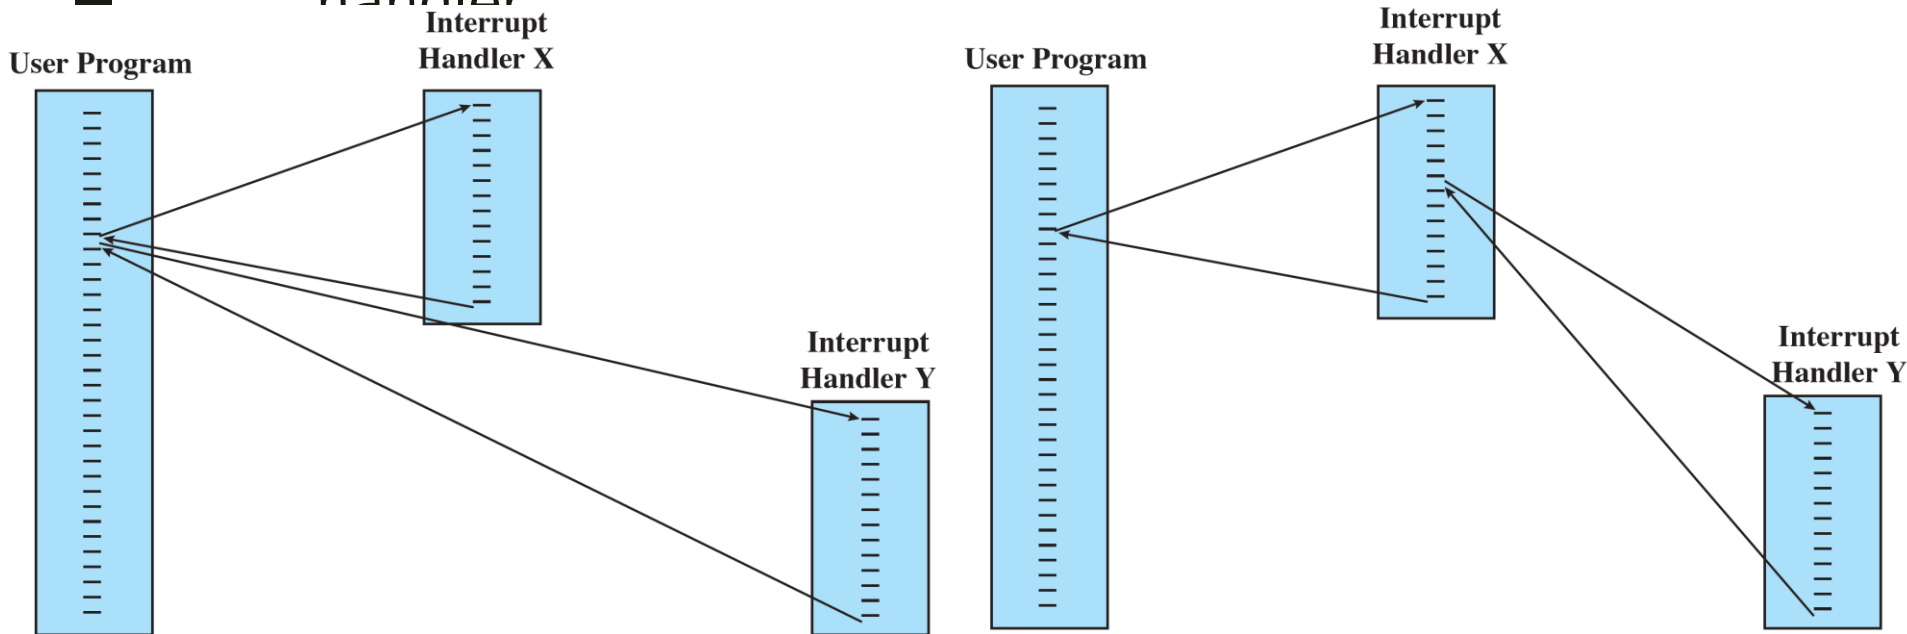
1. *Stores the rest of the general-purpose registers being used by the interrupted process*
2. *Does the rest of interrupt handling operation*

When Handler Ends

- Software does the following:
 - *Handler code restores the saved general registers*
- Hardware does the following:
 - *Restores PC, SP, EFLAGS from Kernel Stack*
 - *Reenable Interrupts*
 - *Switch to user mode*

Sequential vs Nested Interrupt Handling

- In x86, other interrupts are disabled to avoid confusion
 - *This keeps things simple, however, no levels of priority among interrupts*
- Therefore, many systems support nested interrupt handler



(a) Sequential interrupt processing

(b) Nested interrupt processing

Summary: User/Kernel (Privileged) Mode

